


```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [3]: df=pd.read_csv(r"C:\Users\manis\OneDrive\Desktop\ParisHousing.csv")
df
```

Out[3]:

	squareMeters	numberOfRooms	hasYard	hasPool	floors	cityCode	cityPartRang
0	75523	3	0	1	63	9373	
1	80771	39	1	1	98	39381	
2	55712	58	0	1	19	34457	
3	32316	47	0	0	6	27939	1
4	70429	19	1	1	90	38045	
...	...	...	...	...	...	...	
9995	1726	89	0	1	5	73133	
9996	44403	29	1	1	12	34606	
9997	83841	3	0	0	69	80933	1
9998	59036	70	0	0	96	55856	
9999	1440	84	0	0	49	18412	

10000 rows × 20 columns



```
In [4]: df=df[['squareMeters','price']]
df.columns=['sm','pr']
```


```
In [5]: df.head(10)
```

```
Out[5]:
```

	sm	pr
0	75523	7559081.5
1	80771	8085989.5
2	55712	5574642.1
3	32316	3232561.2
4	70429	7055052.0
5	39223	3926647.2
6	58682	5876376.5
7	86929	8696869.3
8	51522	5154055.2
9	39686	3970892.1

```
In [6]: sns.lmplot(x='sm',y='pr',data=df,order=2,ci=None)
```

```
Out[6]: <seaborn.axisgrid.FacetGrid at 0x205f3960ed0>
```



```
In [7]: df.describe()
```

```
Out[7]:
```

	sm	pr
<b>count</b>	10000.00000	1.000000e+04
<b>mean</b>	49870.13120	4.993448e+06
<b>std</b>	28774.37535	2.877424e+06
<b>min</b>	89.00000	1.031350e+04
<b>25%</b>	25098.50000	2.516402e+06
<b>50%</b>	50105.50000	5.016180e+06
<b>75%</b>	74609.75000	7.469092e+06
<b>max</b>	99999.00000	1.000677e+07

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    sm      10000 non-null    int64
1    pr      10000 non-null    float64
dtypes: float64(1), int64(1)
memory usage: 156.4 KB
```

```
In [9]: df.fillna(method='ffill',inplace=True)
```

C:\Users\manis\AppData\Local\Temp\ipykernel\_16088\4116506308.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method='ffill',inplace=True)
```

```
In [11]: X=np.array(df['sm']).reshape(-1,1)
y=np.array(df['pr']).reshape(-1,1)
df.dropna(inplace=True)
```

C:\Users\manis\AppData\Local\Temp\ipykernel\_16088\2340689882.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
df.dropna(inplace=True)
```

```
In [12]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

```
0.9999987468023971
```


```
In [13]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='b')
```

```
plt.plot(X_test,y_pred,color='k')  
plt.show()
```



```
In [14]: df500=df[:][:500]  
sns.lmplot(x='sm',y='pr',data=df500,order=1,ci=None)
```


```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x2058a501790>
```



```
In [15]: df500.fillna(method='ffill',inplace=True)
X=np.array(df500['sm']).reshape(-1,1)
y=np.array(df500['pr']).reshape(-1,1)
df500.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print("Regression: ",regr.score(X_test,y_test))
```

Regression: 0.9999986053702787

```
In [16]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```



```
In [17]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score: ",r2)
```

R2 score: 0.9999986053702787

In [ ]: